

Ragionato per un paio di secondi

## Z-Ant IR & Zig Renderer – Architecture Deep Dive

(A practical orientation for anyone who will add new UOp-lowerings or renderer back-ends)

### 1. Big-Picture Flow

ONNX → High-level Z-Ant passes (operator lowering) → UOpBuilder(slice) (SSA micro-IR) → Zig renderer (pretty prints) →

Stage	Responsibility	Key files
Lower Xxx passes	Break one ONNX node into a handful of <i>UOps</i> (lowering).	lower_add.zig , lower_reduce_mean.zig , ...
UOpBuilder	Owens the growing SSA slice. Guarantees <i>single assignment</i> . Dupes all owned memory.	ir.zig section 4
Renderer (Zig)	Walks the slice once and prints valid Zig code that performs the same scalar work.	zig_renderer.zig + renderers per opcode
Runtime	We just <code>zig build</code> the produced code and call <code>generated_kernel()</code> from tests.	unit tests

### 2. The IR: uop , UOpType , Any , DType

#### 2.1 UOpType

A closed enum of *all* primitive operations that any backend must understand. *They are intentionally tiny* – think of them as the “assembly” of tensor code.

```
DEFINE_GLOBAL // top-level tensor
VIEW          // alias + broadcast info
GEP           // pointer arith (GetElementPtr)
LOAD / STORE  // scalar mem-ops
RANGE / ENDRANGE // counted loops
ADD, MUL, NEG... // scalar ALU
```

```
MULACC          // fused FMA into DEFINE_ACC
...
```

## 2.2 UOp

```
pub const UOp = struct{
    id   : usize,          // index in the slice (SSA id == pos)
    op   : UOpType,        // opcode tag
    dtype: DType,          // element type of *result*
    src  : []const usize,  // producer ids (already defined ops)
    arg  : ?Any,           // optional payload
};
```

- **SSA guarantee:** every ID is defined **once**; later ops only read it.
- `src` is always duplicated by the builder → you own a private slice.

## 2.3 Any – payload

Single-slot tagged union that carries the odd bits of metadata:

Variant	Used by	Example
int / float	CONST	literal “5”, “0.5f”
loop_bounds	RANGE	{start=0,end=128}
mem_info	GEP	{ base=<id>, offset, stride }
view_meta	VIEW	{ shape=[2,3], strides=[3,1] }

Add more when a new primitive needs it.

## 3. Lowering helpers – UOpBuilder

```
const id = b.push(.ADD, .f32, &.{lhs, rhs}, null);
```

- `push()` duplicates **both** `src` and (if needed) the slices stored inside `Any.view_meta`.
- `toOwnedSlice()` hands ownership of the finished program to the caller.
- `deinit()` frees *everything* (each `src`, each duplicated `shape/stride`).

**Rule of thumb:** If you allocate memory inside `Any`, duplicate it before storing and free it in `deinit()`.

## 4. Renderer architecture

```

zig_renderer.zig
├─ identify_buffers() // builds BufferInfo map
├─ render_as_function() // top-level entry
│   1. emits Zig fn signature
│   2. allocates buffers (DEFINE_GLOBAL)
│   3. creates ptr_map (id → var name)
│   4. walks slice -> render_uop()
│   5. returns output slice
├─ render_uop() // one big switch
│   ├─ MemoryRender.render() // LOAD / STORE / CONST / DEFINE_GLOBAL
│   ├─ GepRender.render() // pointer math
│   ├─ ArithmeticRender.render() // ADD/MUL/...
│   ├─ UnaryRender.render() // NEG/CAST/EXP2...
│   ├─ ControlFlowRender // RANGE / ENDRANGE
│   ├─ manage VIEW meta // adds to view_map, no code
│   └─ DEFINE_ACC / MULACC // inline helpers

```

### 4.1 Maps kept during rendering

Name	Type	Purpose
buffer_map	HashMap(usize, BufferInfo)	semantic info for each DEFINE_GLOBAL (shape, name, is_input)
view_map	HashMap(usize, ViewInfo)	stores {shape, strides} for each VIEW id
ptr_map	HashMap(usize, []const u8)	final Zig <i>identifier</i> (variable name) that holds the value/pointer for an SSA id
rendered_ids	HashSet(usize)	ensures we print each op once

### 4.2 Naming convention (auto-generated)

Prefix	What it is
input_	function parameter slice
output_	final output buffer slice
addr_	<b>usize</b> holding a calculated address ( GEP )
buf_	scalar temporaries ( ADD result, etc.)

Prefix	What it is
idx_	loop induction variable ( RANGE )
acc_	accumulator (DEFINE_ACC)
view_	alias id (no code emitted)

## 5. GepRender – deep dive

Goals: turn a high-level GEP into

```
const addr_7 = @intFromPtr(base.ptr) + (offset_expr) * @sizeof(f32);
```

Steps:

### 1. Base pointer selection

- If `base_id` is a `DEFINE_GLOBAL` slice that is `input_*` or `output_*` → need `.ptr` because slices are `{ ptr, len }`.
- If it is an internal pointer ( `addr_*` , `acc_*` ) → already a raw pointer.

**2. Offset expression** *If there is a VIEW we respect its per-axis stride and optional broadcast (stride == 0). 1-D index form: supports rank-1 and rank-2 by unflattening. Full index form: just  $\sum (\text{idx} \times \text{stride})$  skipping broadcast axes. Raw buffer (no VIEW) assumes plain row-major layout (rank-1 or rank-2).*

**3. Emit final line** – multiply offset by `@sizeof(dtype)` .

Helpers:

```
castIndex()           // "idx_i32" → "@as(usize,@intCast(idx_i32))"
emitTerm()            // prints "+ (expr*stride)" skipping stride==0
ArenaAllocator        // every temporary string is arena-allocated ⇒ freed at end
```

## 6. MemoryRender – LOAD / STORE / CONST

- **CONST** → `var buf_4: f32 = 0.5; _ = &buf_4;`
- **LOAD** → read through the calculated pointer, result bound to `var buf_9` .
- **STORE** → write scalar value *directly* via pointer cast.

All three look up variable names in `ptr_map` .

## 7. Control-flow ( `RANGE` / `ENDRANGE` )

---

```
var idx_3: i32 = 0; // RANGE
while (idx_3 < end) : (idx_3 += 1) {
    ...
} // ENDRANGE
```

Indentation depth is tracked in `loop_indent` so nested loops indent correctly.

## 8. Adding a new lowering pass

---

### 1. In `LowerXxx.zig`

- Calculate output `shape` , `strides` , etc.
- Emit VIEWS / GEPs / ALU UOps through `uopBuilder` .

### 2. Unit-test

- Dump the slice ( `uop.dump` ) to make sure the sequence is valid.

### 3. Rendering

- If your new op re-uses existing primitives (often the case), no renderer work is needed.
- Otherwise add a small renderer file: follow the pattern of `arithmetic_render.zig` .

### 4. Update `render_uop` **switch** to call your renderer.

## 9. Adding a brand-new primitive

---

1. Extend `uopType` , add a case.
2. Decide if it needs metadata → add a field to the `Any` union.
3. Teach `uopBuilder.push()` to duplicate that field when it appears.
4. Write a renderer and hook it in `render_uop` .
5. Write one minimal lowering + test so CI exercises it.

## 10. Memory ownership rules

Who allocates	What	Who frees
Lowering pass	slices in <code>Any</code>	<b>UOpBuilder.deinit</b>
Renderer	<code>acc_*</code> , <code>view_*</code> names via <code>allocPrint</code>	Freed at end of <code>render_as_function</code> (only those prefixes)
<code>identify_buffers</code>	<code>BufferInfo.name</code> , <code>.shape</code>	Freed in <code>ZigRenderer.deinit</code>

**Never** store a pointer into caller-owned memory inside a `uop` without duplicating it first.

## 11. Common gotchas

- Forgetting to append `.ptr` for slice parameters  $\Rightarrow$  ‘expected pointer, found `[]const f32`’.
- VIEW rank mismatch: `uop.src.len-1` must equal `view_meta.strides.len` in the multi-index form.
- Leaked arena strings: only allocate with the dedicated arena inside the renderer; it is `deinit()` ed every GEP invocation.

## 12. Checklist for new contributors

1. **Lowering** produces valid SSA? run `dump()` and eyeball ordering.
2. **Slice safety**: every `shape[]`, `strides[]` duplicated inside `Any`?
3. **Renderer**: variable names go through `ptr_map`, never hard-code `id`  $\rightarrow$  `name`.
4. **Unit tests**:
  - `zig build test` passes without “leaked” or “expected pointer” errors.
  - Dumped Zig shows correct address arithmetic.
5. **Documentation**: update this file if you add new primitives.

With these conventions in mind, extending Z-Ant either on the *front* (lowerings) or the *back* (renderers) should be a predictable and safe process. Happy hacking!